



# Exploratory Data Analysis

# What does the data look like?



<dataset>



gather - (Contains system logs ,located at 'gather/<host\_name>/logs/')



labels - Contains the ground truth of the dataset that indicates which events are related to attacks.



rules - contains the label rules



processing - Contains the source code that was used to generate the labels.



environment - Contains the source code that was used to deploy the testbed and run the simulation

dataset.yml - specifies the start and end time of the simulation.

# Data Engineering

```
type=USER_AUTH msg=audit(1642999060.603:2226): pid=27950 uid=33 auid=4294967295
ses=4294967295 msg='op=PAM:authentication acct="jhall" exe="/bin/su" hostname=? addr=?
terminal=/dev/pts/1 res=success' type=USER_ACCT msg=audit(1642999060.603:2227): pid=27950
uid=33 auid=4294967295 ses=4294967295 msg='op=PAM:accounting acct="jhall" exe="/bin/su"
hostname=? addr=? terminal=/dev/pts/1 res=success' type=CRED_ACQ
msg=audit(1642999060.615:2228): pid=27950 uid=33 auid=4294967295 ses=4294967295
msg='op=PAM:setcred acct="jhall" exe="/bin/su" hostname=? addr=? terminal=/dev/pts/1 res=success'
type=USER_START msg=audit(1642999060.627:2229): pid=27950 uid=33 auid=4294967295
ses=4294967295 msg='op=PAM:session_open acct="jhall" exe="/bin/su" hostname=? addr=?
terminal=/dev/pts/1 res=success'
```

```
{"line": 1860, "labels": ["attacker_change_user", "escalate"], "rules": {"attacker_change_user":
["attacker.escalate.audit.su.login"], "escalate": ["attacker.escalate.audit.su.login"]}} {"line": 1861, "labels":
["attacker_change_user", "escalate"], "rules": {"attacker_change_user":
["attacker.escalate.audit.su.login"], "escalate": ["attacker.escalate.audit.su.login"]}} {"line": 1862, "labels":
["attacker_change_user", "escalate"], "rules": {"attacker_change_user":
["attacker.escalate.audit.su.login"], "escalate": ["attacker.escalate.audit.su.login"]}} {"line": 1863, "labels":
["attacker_change_user", "escalate"], "rules": {"attacker_change_user":
["attacker.escalate.audit.su.login"], "escalate": ["attacker.escalate.audit.su.login"]}}
```



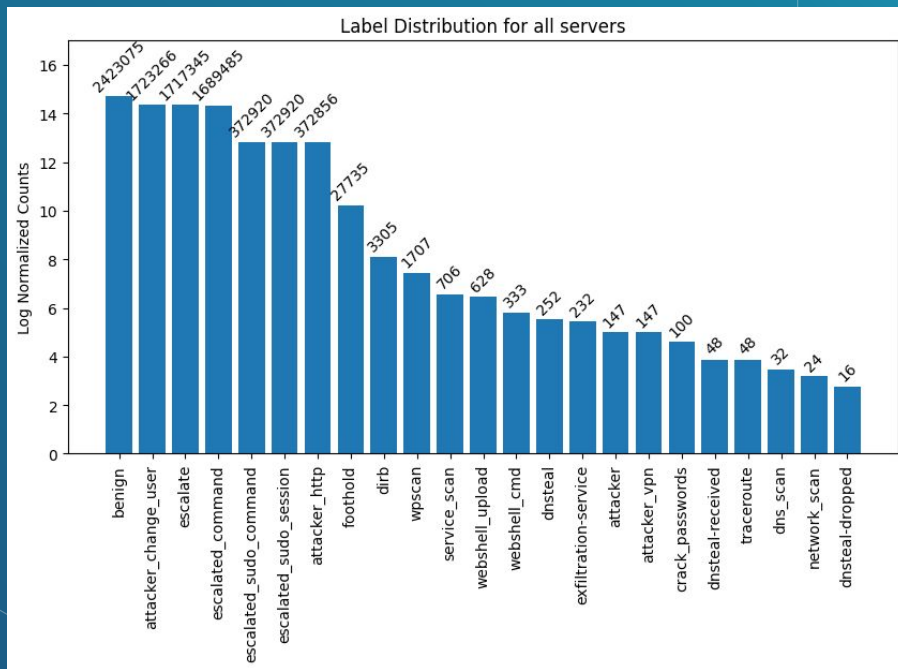
# Data Engineering

|        | line   | log_line   | labels | rules  | gather_path                                    | labels_path                                    | server_name |
|--------|--------|--|--------|--------|--|--|-------------|
| 0      | 1      | Jan 23 06:25:12 intranet-server CRON[27138]: p...  | benign | benign | wardbeck/gather/intranet_server/logs/auth.log  | wardbeck/labels/intranet_server/logs/auth.log  | wardbeck    |
| 1      | 2      | Jan 23 06:39:01 intranet-server CRON[27297]: p...  | benign | benign | wardbeck/gather/intranet_server/logs/auth.log  | wardbeck/labels/intranet_server/logs/auth.log  | wardbeck    |
| 2      | 3      | Jan 23 06:39:01 intranet-server CRON[27297]: p...  | benign | benign | wardbeck/gather/intranet_server/logs/auth.log  | wardbeck/labels/intranet_server/logs/auth.log  | wardbeck    |
| 3      | 4      | Jan 23 06:47:01 intranet-server CRON[27367]: p...  | benign | benign | wardbeck/gather/intranet_server/logs/auth.log  | wardbeck/labels/intranet_server/logs/auth.log  | wardbeck    |
| 4      | 5      | Jan 23 06:47:05 intranet-server CRON[27367]: p...  | benign | benign | wardbeck/gather/intranet_server/logs/auth.log  | wardbeck/labels/intranet_server/logs/auth.log  | wardbeck    |
| ...    | ...    | ...  | ...    | ...    | ...  | ...  | ...         |
| 329762 | 299051 | Jan 23 23:47:10 dnsmasq[14755]: reply motd.ubu...  | benign | benign | wardbeck/gather/inet-firewall/logs/dnsmasq.log | wardbeck/labels/inet-firewall/logs/dnsmasq.log | wardbeck    |
| 329763 | 299052 | Jan 23 23:57:15 dnsmasq[14755]: query[A AAA] in... | benign | benign | wardbeck/gather/inet-firewall/logs/dnsmasq.log | wardbeck/labels/inet-firewall/logs/dnsmasq.log | wardbeck    |
| 329764 | 299053 | Jan 23 23:57:15 dnsmasq[14755]: forwarded intr...  | benign | benign | wardbeck/gather/inet-firewall/logs/dnsmasq.log | wardbeck/labels/inet-firewall/logs/dnsmasq.log | wardbeck    |
| 329765 | 299054 | Jan 23 23:57:15 dnsmasq[14755]: nameserver 127...  | benign | benign | wardbeck/gather/inet-firewall/logs/dnsmasq.log | wardbeck/labels/inet-firewall/logs/dnsmasq.log | wardbeck    |
| 329766 | 299055 | Jan 23 23:57:15 dnsmasq[14755]: reply intranet...  | benign | benign | wardbeck/gather/inet-firewall/logs/dnsmasq.log | wardbeck/labels/inet-firewall/logs/dnsmasq.log | wardbeck    |

329767 rows x 7 columns



# Initial Label Distribution





# Rules

- Rules folder contains the description of each labels and the pattern behind each attack that how it was classified

| type                       | id                                       | labels                        | description   |
|----------------------------|--|-------------------------------|---|
| elasticsearch.sequence     | attacker.foothold.apache.access          | ['attacker_http', 'foothold'] | This rule matches the attackers recorded HTTP traffic to access log lines based on the web paths. Note that we prefix the web paths with the servers FQDN in the parsing phase to match the PCAP records.   |
| elasticsearch.sequence     | attacker.foothold.pcap.requests          | ['attacker_http', 'foothold'] | This rule matches attacker HTTP request pcap logs with their HTTP responses when the response has already been marked as attacker traffic. We do this so we can apply followup rules only on requests that have not been marked yet. Such cases can happen when the response did not make it to the attacker.   |
| elasticsearch.sub_query    | attacker.foothold.apache.access_dropped  | ['attacker_http', 'foothold'] | This rule tries to match attacker requests that we where unable to match to a labeled response with access log entries. Such cases can happen if the corresponding response gets lost in the network or otherwise is not sent.  |
| elasticsearch.sequence     | attacker.foothold.apache.error           | ['attacker_http', 'foothold'] | This rule matches the attackers recorded HTTP traffic to error log lines based on the web server file the error occurred for. We do this by converting the reported server file paths to web paths in the parsing phase.  |
| elasticsearch.sequence     | attacker.foothold.apache.access_error    | ['attacker_http', 'foothold'] | This rule matches already identified access lines to error log lines.   |
| elasticsearch.sub_query    | attacker.foothold.apache.error_substring | ['attacker_http', 'foothold'] | Apache error logs sometimes do not have an url.full attribute, especially events related to the WPScan. Therefore they cannot be labeled by grouping access and error logs with the same url.full value. This rule tries to resolve this issue by checking whether the url occurs in other attributes where it could appear depending on the type of the error event. Note that url.original is used since only this part sometimes occurs rather than the full url. Also note that both regexp and match_phrase is used as elasticsearch sometimes analyzes fields which prevents simple regex matching. |
| elasticsearch.sub_query    | attacker.foothold.apache.error_index     | ['attacker_http', 'foothold'] | This rule applies the attacker http label to all errors produced by requests to directories that map to an index.php file. This is done by first getting all labeled directory requests and then searching for matching error lines with index.php prefixed.  |
| elasticsearch.parent_query | attacker.foothold.apache.error_access    | ['attacker_http', 'foothold'] | This rule looks for unlabeled error messages resulting from VPN server traffic within the attack time and tries to match it to an already labeled access log row.   |
| elasticsearch.parent_query | attacker.foothold.apache.php_warnings    | ['attacker_http', 'foothold'] | This rule tries to apply missing labels to error lines that result from multiple PHP errors/warnings in a single request. First all PHP errors that are potentially from the attacker are retrieved and then we try to find a labeled parent by matching pid, port, url.full and the approximate error time.  |
| elasticsearch.query        | attacker.dirb.time                       | ['dirb']                      | This labels attacker http traffic within the recorded dirb execution time with dirb.  |
| elasticsearch.query        | attacker.wpscan.time                     | ['wpscan']                    | This labels attacker http traffic within the recorded wpscan execution time with wpscan.  |

# Problem with Rules Table

| type                       | id  |
|----------------------------|---|
| elasticsearch.sequence     | attacker.foothold.apache.access             |
| elasticsearch.sequence     | attacker.foothold.pcap.requests             |
| elasticsearch.sub_query    | attacker.foothold.apache.access_dropped     |
| elasticsearch.sequence     | attacker.foothold.apache.error              |
| elasticsearch.sequence     | attacker.foothold.apache.access_error       |
| elasticsearch.sub_query    | attacker.foothold.apache.error_substring    |
| elasticsearch.sub_query    | attacker.foothold.apache.error_index        |
| elasticsearch.parent_query | attacker.foothold.apache.error_access       |
| elasticsearch.parent_query | attacker.foothold.apache.php_warns          |
| elasticsearch.query        | attacker.dirb.time                          |
| elasticsearch.query        | attacker.wpscan.time                        |
| elasticsearch.sequence     | attacker.webshell.upload.seq                |
| elasticsearch.query        | attacker.escalate.webshell.cmd.http         |
| elasticsearch.query        | attacker.escalate.webshell.cmd.http_prepare |
| elasticsearch.sequence     | attacker.escalate.su.login                  |
| elasticsearch.sequence     | attacker.escalate.systemd.newsession.before |
| elasticsearch.sequence     | attacker.escalate.systemd.newsession.after  |
| elasticsearch.sequence     | attacker.escalate.systemd.session           |
| elasticsearch.sequence     | attacker.escalate.sudo.command              |
| elasticsearch.sequence     | attacker.escalate.sudo.open                 |
| elasticsearch.sequence     | attacker.escalate.audit.su.login            |
| elasticsearch.sequence     | attacker.escalate.audit.systemd.session     |
| elasticsearch.sequence     | attacker.escalate.audit.sudo.command.start  |
| elasticsearch.sequence     | attacker.escalate.audit.sudo.command.events |
| elasticsearch.sequence     | attacker.escalate.audit.sudo.command.events |
| elasticsearch.query        | dnsteal.domain.match                        |
| elasticsearch.query        | dnsteal.domain.received                     |
| elasticsearch.query        | dnsteal.domain.dropped                      |
| elasticsearch.parent_query | dnsteal.domain.dropped-retry                |
| elasticsearch.query        | attacker.foothold.dnsmasq.wpscan            |
| elasticsearch.query        | attacker.foothold.dnsmasq.webshell          |
| elasticsearch.query        | attacker.foothold.dnsmasq.dns_scan_probe    |

- ◇ 48+ Different Types of Attacks made it harder to categorize and make Large Language Model to train to categorize them

# Solution: Merged Labels

- ◇ Merged a similar attack type to the “merged label” in order to make a categorization easier

| type                       | Merged Label      | id                  | labels                        | description   |
|----------------------------|-------------------|---------------------|-------------------------------|---|
| elasticsearch.sequence     | HTTP Attack       | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule matches the attackers recorded <b>HTTP</b> traffic to     |
| elasticsearch.sequence     |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule matches attacker <b>HTTP</b> request pcap logs with       |
| elasticsearch.sub_query    |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule tries to match attacker requests that we where            |
| elasticsearch.sequence     |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule matches the attackers recorded <b>HTTP</b> traffic to     |
| elasticsearch.sequence     |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule matches already identified access lines to error          |
| elasticsearch.sub_query    |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | Apache error logs sometimes do not have an url.full attrib          |
| elasticsearch.sub_query    |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule applies the attacker <b>http</b> label to all errors prod |
| elasticsearch.parent_query |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule looks for unlabeled error messages resulting fro          |
| elasticsearch.parent_query |                   | attacker.foothold.. | ['attacker_http', 'foothold'] | This rule tries to apply missing labels to error lines that re      |
| elasticsearch.query        |                   | attacker.dirb.time  | ['dirb']                      | This labels attacker <b>http traffic</b> within the recorded dirb   |
| elasticsearch.query        |                   | attacker.wpscan.t   | ['wpscan']                    | This labels attacker <b>http traffic</b> within the recorded wpsc   |
| elasticsearch.sequence     | Web Shell Request | attacker.webshell.  | ['webshell_upload']           | This rule labels the web shell upload step by matching th           |
| elasticsearch.query        |                   | attacker.escalate.  | ['webshell_cmd']              | This rule matches the web shell web requests via the rec            |
| elasticsearch.query        |                   | attacker.escalate.  | ['webshell_cmd', 'esc         | This rule matches the web shell web requests via the rec            |
| elasticsearch.sequence     |                   | attacker.escalate.  | ['attacker_change_us          | This rule labels auth log rows resulting from the attacker          |

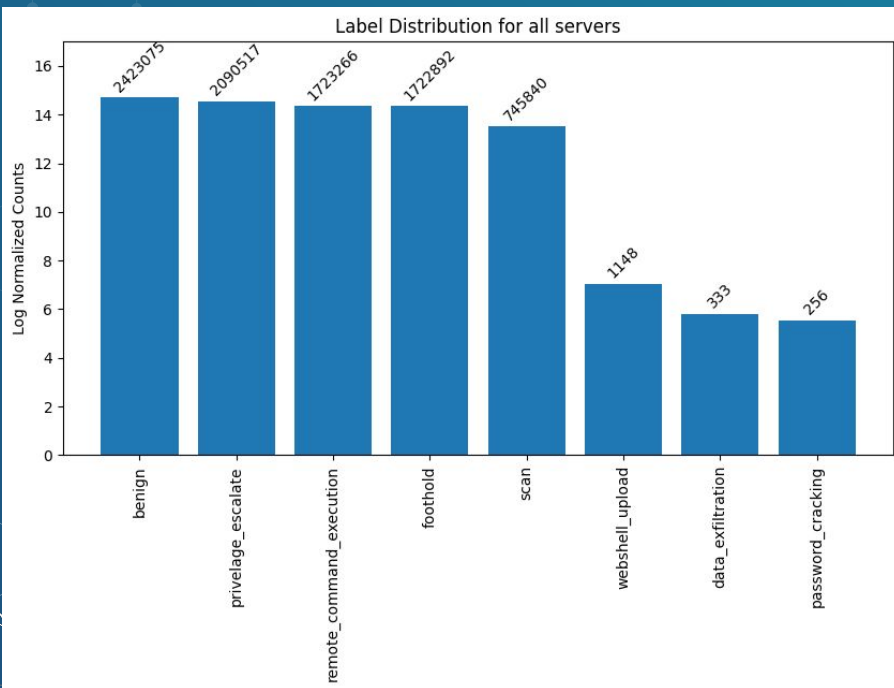


# Reduce Labels by Grouping

48+ →

- ◇ Benign
- ◇ Privilege Escalation
- ◇ Remove Command Execution
- ◇ Foothold
- ◇ Webshell Upload
- ◇ Data Exfiltration
- ◇ Password Cracking

# Refined Label Distribution



$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

$$\text{Accuracy} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives} + \text{True Negatives} + \text{False Negatives}}$$

$$\text{F1 Score} = 2 * \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

# Initial N-Gram

```
Top 100 3-grams:
('','',''): 51840 occurrences
('','',''): 51840 occurrences
('','',''): 32640 occurrences
({'','','pct'): 32640 occurrences
('','','pct'): 32640 occurrences
('','',''): 26880 occurrences
('','',''): 17280 occurrences
('norm','',''): 17280 occurrences
('','','norm'): 13440 occurrences
('','','norm'): 13440 occurrences
({'','',''): 11928 occurrences
('','',''): 9600 occurrences
('pct','',''): 9084 occurrences
('','','system'): 7680 occurrences
('','','intranet-server'): 5760 occurrences
('','','intranet-server'): 5760 occurrences
('','','name'): 5760 occurrences
('','','name'): 5760 occurrences
('name','',''): 5760 occurrences
('version','',''): 5760 occurrences
('','',''): 5557 occurrences
({'','',''): 5352 occurrences
('intranet-server','',''): 3840 occurrences
('type','',''): 3840 occurrences
('','','system'): 3840 occurrences
('','','@'): 3840 occurrences
({'','','cpu'): 3840 occurrences
('','','cpu'): 3840 occurrences
('cpu','',''): 3840 occurrences
('','','system'): 3840 occurrences
('system','',''): 3840 occurrences
({'','','norm'): 3840 occurrences
```

- ◇ N Gram was initially used to find the most frequency of pattern, but it didn't return meaningful data since it's the log file that it returned some format instead of the pattern



# Advanced N-Gram

◇ Bronze Live Code Demonstration



The background of the slide is a solid blue color with a subtle, repeating pattern of white hexagons. Some of these hexagons are interconnected by thin white lines, creating a network-like or molecular structure. In the top-left corner, there is a small, solid blue hexagon.

# Discussion

- ◇ Because the data is so unstructured, it is hard to perform meaningful classical EDA
- ◇ Data cleaning and Modern NLP techniques such as LLM-tokenization allowed us to gain more meaningful insights





# Questions?